

## MACHINE LEARNING AND BIG DATA – ASSIGNMENT 2 – MULTIVARIABLE REGRESSION

Multivariable regression is actually pretty similar to linear regression. The only difference from that is we have “multiple variables (parameters)”.

Because of that, we will now have a model with multiple parameters

### Model

Previously:  $f_{w,b}(x) = wx + b$

$$f_{w,b}(x) = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$

example:

$$f_{w,b}(x) = 0.1x_1 + 4x_2 + 10x_3 + -2x_4 + 80$$


size                      years  
#bedrooms              #floors

$$f_{w,b}(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n + b$$

multiple linear regression  
(not multivariate regression)

Here is how each of our variables are respect to the target



I will explain my functions one by one:

- Z-score
- Compute cost
- Compute descent
- Compute gradient descent

## 1) Z-score

Here, we will use this z-score function to “normalise” our X inputs

When constructing our model, we will normalise X values (which we got from houses.txt file) and use normalised X values

```
def zscore_normalize_features(X):  
    """  
    computes X, zscore normalized by column  
  
    Args:  
        X (ndarray (m,n)) : input data, m examples, n features  
  
    Returns:  
        X_norm (ndarray (m,n)): input normalized by column  
        mu (ndarray (n,)) : mean of each feature  
        sigma (ndarray (n,)) : standard deviation of each feature  
    """  
  
    mu = np.mean(X, axis=0) # Calculate the mean of each feature  
    sigma = np.std(X, axis=0) # Calculate the standard deviation of each feature  
    X_norm = (X - mu) / sigma # Normalize the input data by subtracting the mean and dividing by the standard deviation  
  
    return (X_norm, mu, sigma)
```

A SMALL BUT STILL AN IMPORTANT THING TO POINT OUT is that, we will also use the “sigma and mu” values in smaller sample examples...

## 2) Compute Cost

Compute cost, calculates the error... when we say "error" we mean "the 'general' rate how far our 'estimations' are from the 'target' "

$$J(\mathbf{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

$$f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = \mathbf{w} \cdot \mathbf{x}^{(i)} + b$$

Here is my compute cost function...

```
def compute_cost(X, y, w, b):  
    """  
    compute cost  
    Args:  
        X (ndarray (m,n)): Data, m examples with n features  
        y (ndarray (m,)) : target values  
        w (ndarray (n,)) : model parameters  
        b (scalar)       : model parameter  
    Returns  
        cost (scalar)    : cost  
    """  
  
    m = len(y) # Number of examples  
  
    # Compute predictions  
    predictions = np.dot(X, w) + b # X * w + b  
  
    # Compute squared error  
    squared_error = np.square(predictions - y)  
  
    # Compute mean squared error  
    cost = 1 / (2 * m) * np.sum(squared_error)  
  
    return cost
```

### 3) Compute gradient

Here, we are taking the derivative of the cost function (cost function is coming from compute\_cost)

However, it is a little bit different to “finding derivative in multivariable” than linear (single variable function)

Since we have multiple variables, “respect to which one” will we be taking the derivative?

➔ Answer is, respect to all...

- Taking derivative of each variable is taking “partial derivative”s of cost function
- If we take the derivative of all variables (if we find all partial derivatives) (after summing them) we will have the “derivative of cost function”

$$\frac{\partial J(\mathbf{w}, b)}{\partial w_j} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$
$$\frac{\partial J(\mathbf{w}, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)})$$

```
def compute_gradient(X, y, w, b):
    """
    Computes the gradient for linear regression
    Args:
        X : (ndarray Shape (m,n)) matrix of examples
        y : (ndarray Shape (m,)) target value of each
        w : (ndarray Shape (n,)) parameters of the model
        b : (scalar) parameter of the model
    Returns:
        dj_dw : (ndarray Shape (n,)) The gradient of the cost function with respect to w
        dj_db : (scalar) The gradient of the cost function with respect to b
    """

    m = len(y)

    predictions = np.zeros(1)
    ...

    print(X)
    print(y)
    print(w)
    print(b)
    ...

    predictions = np.dot(X, w) + b # X * w + b

    error = predictions - y

    # Compute mean squared error
    dj_dw = 1 / (m) * np.dot(X.T, error)

    dj_db = 1 / (m) * np.sum(error)

    #print("dj_db: ", dj_db)

    #print("sss")

    return dj_db, dj_dw
```

#### 4) Compute gradient descent

Gradient descent is same as linear regression... We just need too apply compute\_gradient multiple times until we reach the best w and b values (lowest cost rate)

```
def gradient_descent(X, y, w_in, b_in, cost_function,
                    gradient_function, alpha, num_iters):
    """
    Performs batch gradient descent to learn theta. Updates
    num_iters gradient steps with learning rate alpha

    Args:
        X : (array_like Shape (m,n)) matrix of examples
        y : (array_like Shape (m,)) target value of each e
        w_in : (array_like Shape (n,)) Initial values of para
        b_in : (scalar) Initial value of param
        cost_function: function to compute cost
        gradient_function: function to compute the gradient
        alpha : (float) Learning rate
        num_iters : (int) number of iterations to run gradien

    Returns
        w : (array_like Shape (n,)) Updated values of paramet
            after running gradient descent
        b : (scalar) Updated value of paramete
            after running gradient descent
        J_history : (ndarray): Shape (num_iters,) J at each i
            primarily for graphing later
    """

    # Normalize the features
    X_norm, mu, sigma = zscore_normalize_features(X)

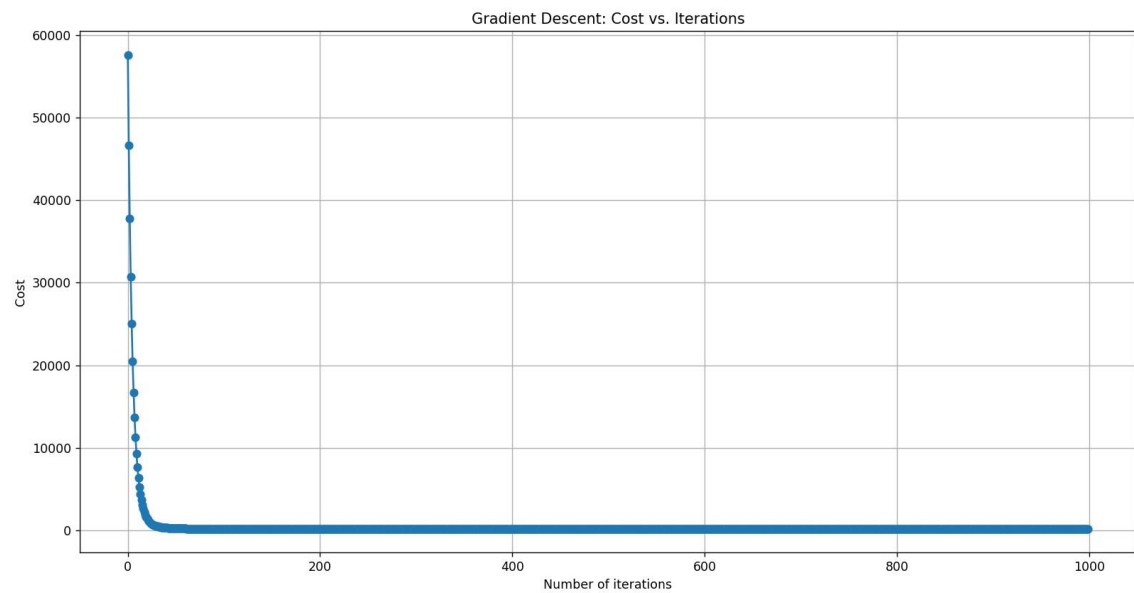
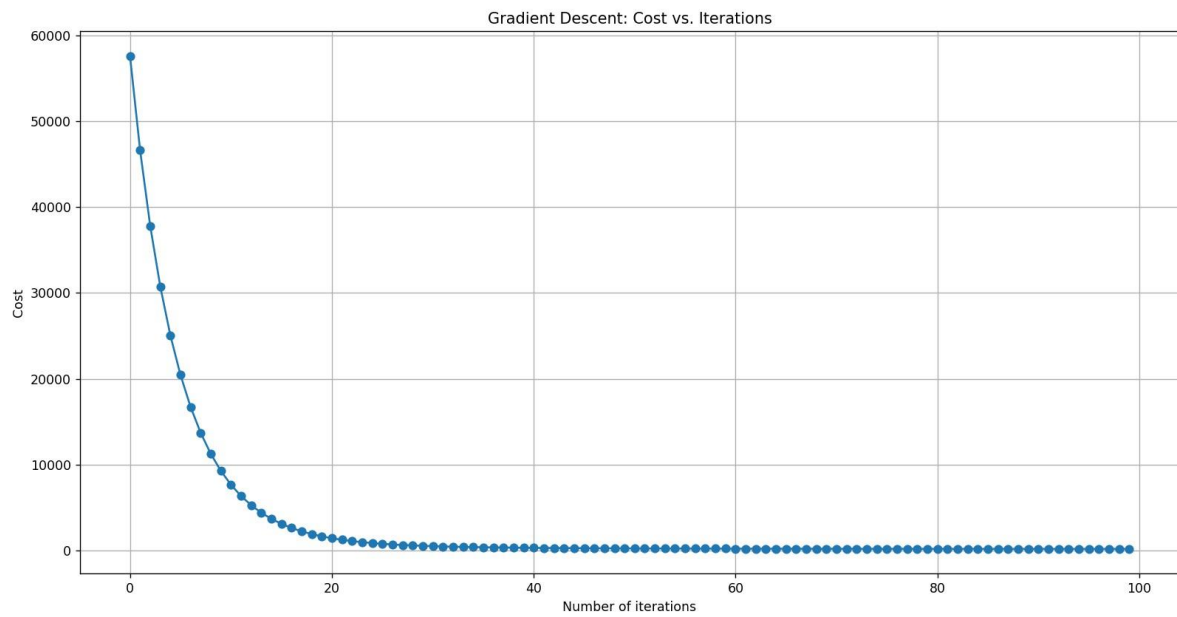
    # Initialize parameters
    w = np.array(w_in)
    b = b_in
    J_history = np.zeros(num_iters)

    for i in range(num_iters):
        # Compute gradients on normalized data
        db, dw = gradient_function(X_norm, y, w, b)


        # Update parameters
        w = w - alpha * dw
        b = b - alpha * db

        # Compute cost on normalized data
        J_history[i] = cost_function(X_norm, y, w, b)

    return w, b, J_history
```



Here is how my model performs after everything



Proof that it passes tests:

```
112 compute_cost_test(compute_cost)
113 compute_gradient_test(compute_gradient)
114
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL	PORTS				
390.8	354.	350.	460.	237.	288.304	282.	249.	304.
332.	351.8	310.	216.96	666.336	330.	480.	330.3	348.
304.	384.	316.	430.4	450.	284.	275.	414.	258.
378.	350.	412.	373.	225.	390.	267.4	464.	174.
340.	430.	440.	216.	329.	388.	390.	356.	257.8 ]

```
W: [110.56039756 -21.26715096 -32.70718139 -37.97015909]
b: 363.15608080808056
Value test: 318.7090923199992
All tests passed!
All tests passed!
PS C:\Users\ardah\UCM>
```



Multi\_linear\_reg.py

```
def compute_gradient(X, y, w, b):
    predictions = np.zeros(1)
    ...
    print(X)
    print(y)
    print(w)
    print(b)
    ...

    predictions = np.dot(X, w) + b # X * w + b

    error = predictions - y

    # Compute mean squared error
    dj_dw = 1 / (m) * np.dot(X.T, error)

    dj_db = 1 / (m) * np.sum(error)

    #print("dj_db: ", dj_db)

    #print("sss")

    return dj_db, dj_dw

def gradient_descent(X, y, w_in, b_in, cost_function,
                    gradient_function, alpha, num_iters):
    """
    Performs batch gradient descent to learn theta. Updates the
    num_iters gradient steps with learning rate alpha

    Args:
        X : (array_like Shape (m,n)) matrix of examples
        y : (array_like Shape (m,)) target value of each example
        w_in : (array_like Shape (n,)) Initial values of parameter
        b_in : (scalar) Initial value of parameter
        cost_function: function to compute cost
        gradient_function: function to compute the gradient
        alpha : (float) Learning rate
        num_iters : (int) number of iterations to run gradient descent

    Returns
        w : (array_like Shape (n,)) Updated values of parameters
            after running gradient descent
        b : (scalar) Updated value of parameter
            after running gradient descent
        J_history : (ndarray): Shape (num_iters,) J at each iteration
            primarily for graphing later

    """
    # Normalize the features
    X_norm, mu, sigma = zscore_normalize_features(X)
    # Initialize parameters
    w = np.array(w_in)
    b = b_in
    J_history = np.zeros(num_iters)

    for i in range(num_iters):
        # Compute gradients on normalized data
        db, dw = gradient_function(X_norm, y, w, b)

        # Update parameters
        w = w - alpha * dw
        b = b - alpha * db

        # Compute cost on normalized data
        J_history[i] = cost_function(X_norm, y, w, b)

    return w, b, J_history
```



